PART TWO

How we get there.

A two-week proof of concept with real data, not a roadmap.

Your deck title here

A one- or two-sentence subhead that clarifies what this deck is about and who it's for.

99%
ACCURACY

500B
PARAMETERS

3×
FASTER

GROUNDX

SUPPLY CHAIN INTELLIGENCE

Your documents are lying to your AI.

Most enterprise RAG systems lose 40% of the signal in their own documents — tables, diagrams, nested context. GroundX ingests them the way a person would read them.

85%
FEWER HALLUCINATIONS

3×
RETRIEVAL PRECISION

$0
IN RE-INGESTION COSTS

THE PROBLEM

Generic RAG misses what matters.

Enterprise documents aren't prose — they're tables, forms, nested captions, and diagrams. A text-first ingestion pipeline flattens all of that and your model retrieves the wrong thing 4 out of 10 queries.

THE ANSWER

Ingest like a person reads.

GroundX parses the structure of the document first — tables as tables, captions bound to figures, sections preserving their hierarchy — then builds retrieval around that structure. The model gets what it asked for.

BUILT FOR THIS

Three capabilities that make enterprise RAG usable.

SCALE

Petabyte-ready ingest.

Parallel pipelines process your entire archive in hours, not weeks. No manual chunking strategy — the system adapts per document class.

SECURITY

Enterprise by default.

SOC2 Type II, HIPAA, and FedRAMP-ready deployments. Your data never leaves your tenant — the ingestion is the integration.

SPEED

Sub-200ms retrieval.

Tuned indices + structural retrieval keep latency low even as the corpus grows. Cached warm paths return in under 80ms at p95.

GROUNDX

HOW IT WORKS

Three steps, two weeks to production.

1

CONNECT

Point GroundX at a document source — S3, SharePoint, an API feed. Ingestion starts immediately; no schema design required.

2

INGEST

Structural parsing runs in parallel across your corpus. Tables, captions, and section hierarchies are preserved end-to-end.

3

RETRIEVE

A single API endpoint returns the structurally-aware chunks your model needs. Drop it into an existing RAG stack in under an hour.

GROUNDX

99%

RETRIEVAL PRECISION ON IL6 WORKLOADS

Independently measured against 12,400 ground-truth queries on the Dell government RAG benchmark. Run the test yourself — we'll ship the dataset.

NEXT STEPS

Let's build this together.

A two-week proof of concept, run on your own corpus, against your own evals. No committee, no six-month integration. Just a clear answer on whether this is worth the rest of your time.

SCHEDULE A WORKING SESSION

READ THE BENCHMARK

GROUNDX

WHAT SHIPS IN THE FIRST 14 DAYS

A working integration, not a slide.

DAY 1–2

Source connector live

S3 / SharePoint / REST ingested against your actual documents.

DAY 3–5

Structural parse validated

Ground-truth compare on 50 representative documents.

DAY 6–10

Retrieval endpoint integrated

Wired into your existing model with your eval harness attached.

DAY 11–14

Precision report delivered

Scored against your queries. Go/no-go decision by end of week 2.

14

DAYS TO FIRST EVAL

0

LINES OF GLUE CODE

$0

IF THE EVAL FAILS